


IronBus segment lifecycle and file cleanup (log-is-WAL): who retires the files, and when



Cleanup ownership (no single actor is blocked by another):
Append actor (#4): seals the active segment, hands immutable index to readers.
Group-commit thread (#6): establishes durable_offset; fatal fsync error freezes the instance.
Checkpointier (#7): publishes the durable watermark (gate 1 for deletion) on roll / 1s / 4 MiB.
Reaper / janitor (#13): retires sealed segments past min(checkpoint, retention, slowest cursor); event-driven on seal, 60s age timer, forced run at 90% disk high-water.
Compactor (#13, optional): re-copies survivor segments, never in-place, single atomic manifest swap.
Backpressure controller (#10): when nothing is eligible and disk fills, durable spills then sheds (drop_new, counted), telemetry drop_oldest, block is opt-in and bounded.

Lessons borrowed: PostgreSQL checkpoints + WAL recycling, RocksDB max_total_wal_size + flush-advances-log, Kafka whole-segment retention (active never deleted), SQLite wal_autocheckpoint, etcd snapshot-then-compact.